

# Apache Spark

Alessandro Margara

[alessandro.margara@polimi.it](mailto:alessandro.margara@polimi.it)

# License

Slides adapted from the official Apache Spark documentation:  
<https://spark.apache.org/docs/latest/index.html>

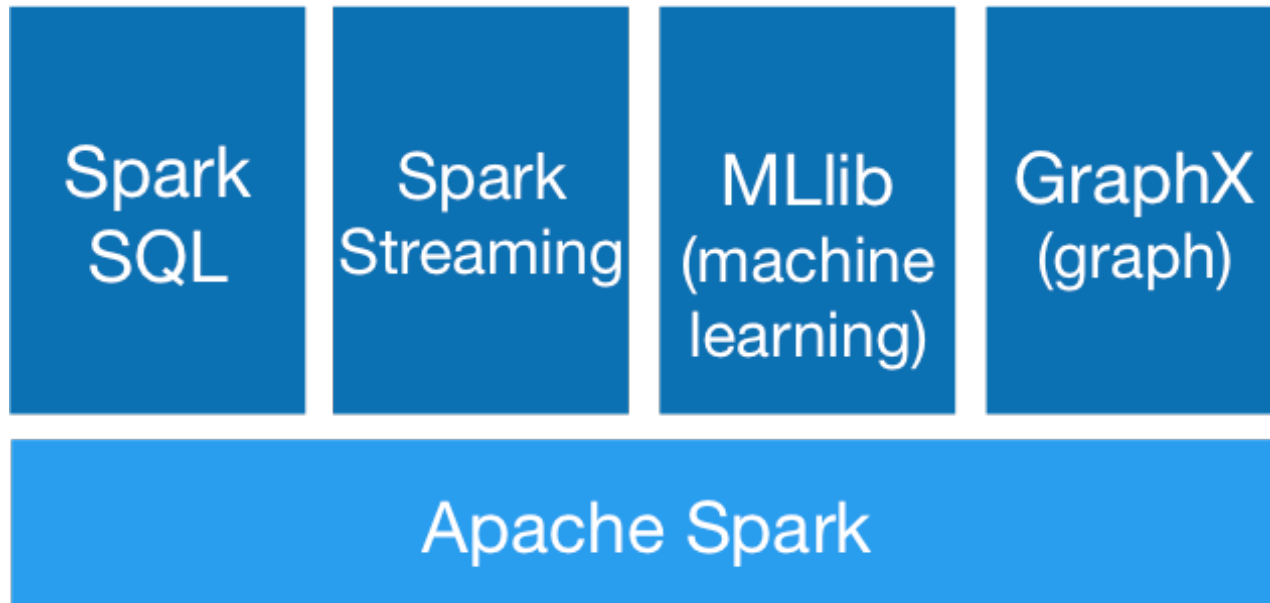
This work is licensed under the Creative Commons Attribution-ShareAlike International Public License



# Overview

- Unified analytics engine for large-scale data processing
  - As we will see, different types of applications: batch, streaming, structured, machine learning, graph processing, ...
- Architecture similar to MapReduce ...
  - Data stored on disk
  - Operators are scheduled on workers
- ... with differences that enable for better performance
  - Data cached in memory
  - Jobs can consist of many “stages”
  - Support for iterative jobs

# A unified analytics engine



# Overview

- A Spark cluster consists of a *master* and one or more workers (*slaves*)
  - Typical configuration: one slave per host, using as many cores as available in the host
- The master
  - Accepts jobs from *driver* programs, and
  - Schedules processing tasks on available slaves

# Overview

- We will write *driver programs* (in Scala) that submit jobs to the cluster
  - Each job consists of various *parallel operations*
- The main abstraction that Spark provides is the RDD (resilient distributed dataset)
  - Collection of elements
  - Partitioned across the nodes of the cluster
  - Can be processed in parallel
  - Fault tolerant

# Initializing Spark

- A Spark program accesses the Spark cluster through a `SparkContext` object
- Contains relevant parameters
  - Name of the Spark application / job
  - Address of the master
  - ...
- Only one context can be active per JVM
  - `stop()` closes a context and enables starting a new one

# Initializing Spark

```
val conf = new SparkConf()  
    .setAppName(appName)  
    .setMaster(master)  
val sc = new SparkContext(conf)
```

...

```
sc.stop()
```



# Initializing spark

- When running in a real cluster
  - The application is packaged in a jar file
  - The jar is submitted to the cluster with a provided script
  - The address of the master is typically not hardcoded, ...
  - ... but extracted from a configuration file
- Local mode
  - Run the driver and the workers in the same JVM
  - For testing and debugging
  - Setting the master to local[n] requests Spark to use n cores to run the workers

# RDDs: initialization

- RDDs can be created from
  - Existing collections in the driver program (a Scala Seq)
  - External datasets
    - Local filesystem
    - HDFS
    - Kafka
    - Several DBMSs
    - ...

# RDDs: initialization

- RDD from local collection

```
val data = Array(1, 2, 3, 4)  
val myRDD = sc.parallelize(data)
```

- RDD from file

```
val myRDD = sc.textFile("filename.txt")
```

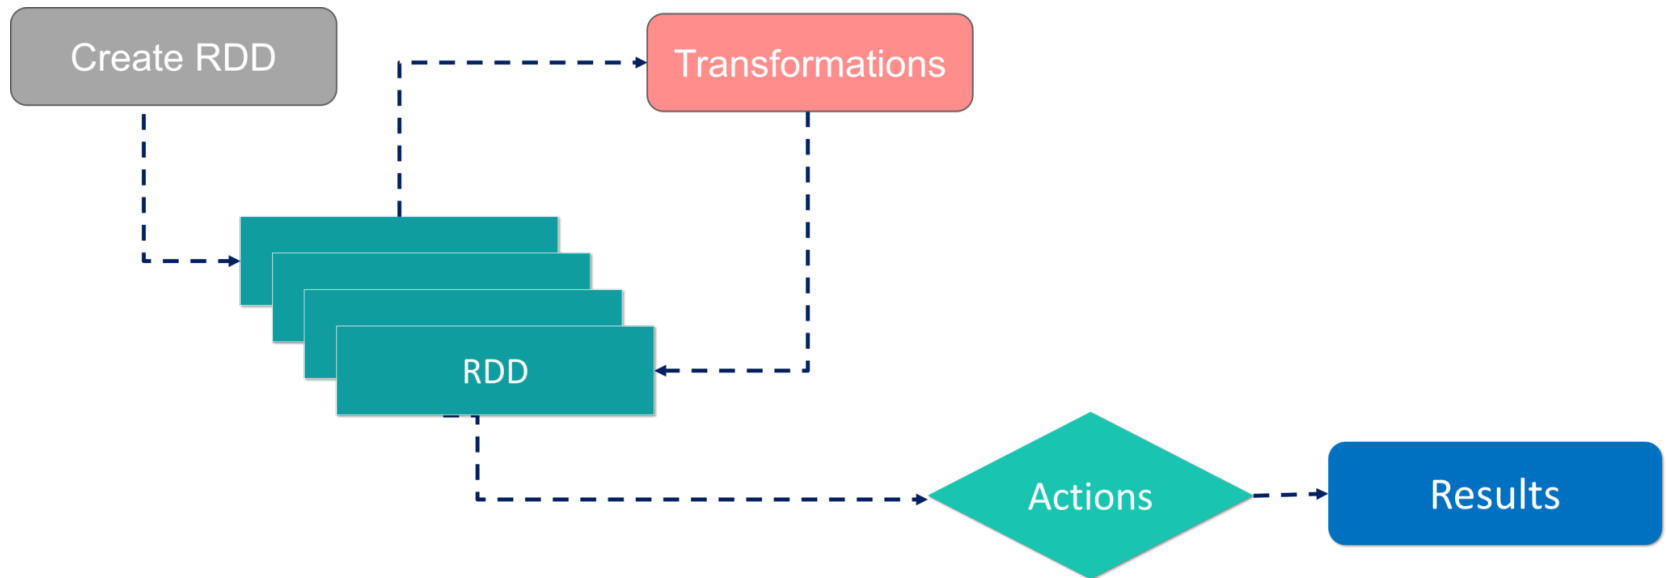
# RDDs: operations

- RDDs support two types of operations
  - *Transformations* create a new RDD from an existing one
  - *Actions* return a value to the driver program after running a computation on the dataset
- All transformations are lazy
  - They do not compute the results when invoked
  - They remember the computation to be implemented, ...
  - ... and execute the computation when an action requires a result to be returned to the driver program

# RDDs: operations

- By default, each transformed RDD is recomputed each time the driver runs (directly or indirectly) an action on it
- The programmer can also persist/cache an RDD in memory (in the workers) for faster access in subsequent queries
- Various storage levels available
  - In-memory objects
  - In-memory serialized objects
  - In-memory + on disk serialized objects
  - On disk serialized objects
  - ...

# RDD: operations



# RDDs: simple example

```
/* Creates an RDD. The number of partitions is decided  
based on the available workers */
```

```
val lines = sc.textFile("data.txt")
```

```
/* Transformation. Map applies a function to each and  
every element of the original RDD. In this case, it  
transforms each string (line) into a number (the length  
of the line) */
```

```
val linesLen = lines.map(s => s.length)
```

```
/* Action. Reduce aggregates all the values in a single  
element and returns the result to the driver. In this  
case, it returns the sum of all the length of all the  
lines) */
```

```
val totLen = linesLen.reduce((a, b) => a + b)
```

# RDDs: simple example

```
val lines = sc.textFile("data.txt")  
val linesLen = lines.map(s => s.length)  
val totLen = linesLen.reduce((a, b) => a + b)
```

- `lines` and `linesLen` are not immediately computed
  - `lines` does not load any data from the file
- When the `reduce` action is invoked, it requests the value of `linesLen`, which requests the value of `lines`
- These values are not persisted
  - Unless the programmer explicitly invokes `cache()` / `persist()`



# RDDs: fault tolerance

- As said, by default RDDs are not persisted
- In the case an RDD is persisted (cached), the cache is fault-tolerant
  - If any partition of an RDD is lost, ...
  - ... it will automatically be recomputed using the transformations that originally created it

# RDDs: some transformations

Transformation	Semantics
map(fun)	Applies fun to each and every element in the source RDD
filter(fun)	Returns a new RDD with all and only the elements e in the source RDD for which fun(e) == true
flatMap(fun)	As map, but fun can return zero, one, or more results for each element in the source RDD

# RDDs: some transformations

Transformation	Semantics
<code>union(otherRDD)</code>	Returns the union of the source RDD and otherRDD
<code>intersection(otherRDD)</code>	Returns the intersection of the source RDD and otherRDD
<code>distinct()</code>	Returns a new RDD that contains the distinct elements in the source RDD

# RDDs: some transformations

Transformation	Semantics
<code>groupByKey()</code>	When called on a RDD of (K, V) pairs, returns a dataset of (K, Iterable<V>) pairs
<code>reduceByKey(fun)</code>	When called on a RDD of (K, V) pairs, returns a new RDD of (K, V) pairs where the values for each key are aggregated using the given reduce function fun (of type (V, V) => V)
<code>aggregateByKey(zero)(seqOp, combOp)</code>	When called on a RDD of (K, V) pairs, returns a RDD of (K, U) pairs where the values of each key are aggregated using the given combine function and a neutral zero value

# RDDs: some transformations

Transformation	Semantics
<code>join(otherRDD)</code>	When called on RDDs of type $(K, V)$ and $(K, W)$ , returns a RDD of $(K, (V, W))$ pairs with all pairs of elements for each key: can be configured for outer joins
<code>cogroup(otherRDD)</code>	When called on RDDs of type $(K, V)$ and $(K, W)$ , returns a RDD of $(K, \text{Iterable}<V>, \text{Iterable}<W>)$ tuples
<code>cartesian(otherRDD)</code>	When called on RDDs of type $T$ and $U$ , returns a RDD of $(T, U)$ pairs (all pairs of elements)

# RDDs: some actions

Transformation	Semantics
<code>reduce(fun)</code>	Aggregate the elements of the source RDD using the function <code>fun</code>
<code>collect()</code>	Returns all the elements of the RDD as an array
<code>count()</code>	Returns the number of elements in the RDD
<code>take(n)</code>	Returns an array with the first <code>n</code> elements in the RDD
<code>saveAsTextFile(path)</code>	Writes the elements as a text file (or set of text files) in the local filesystem or HDFS

# Shuffle operations

- Some operations *shuffle* the data = re-distribute data changing the way they are grouped across partitions
- Shuffle operations involve copying data across workers, making it a complex and costly operation
- Consider for example the classic word count example
  - Data is initially partitioned by document
  - It needs to be re-partitioned by word

# Shuffle operations

- Shuffle is an expensive operation because it involves
  - Serialization/deserialization
  - Disk/network IO
- Internally, shuffle operations can consume memory to store intermediate results while reorganizing the data
  - Data structures kept in memory until they can't fit
  - Then they are spilled to disk
- After the data has been reorganized by key, each key is transmitted to the partition responsible for it



# Spark Architecture

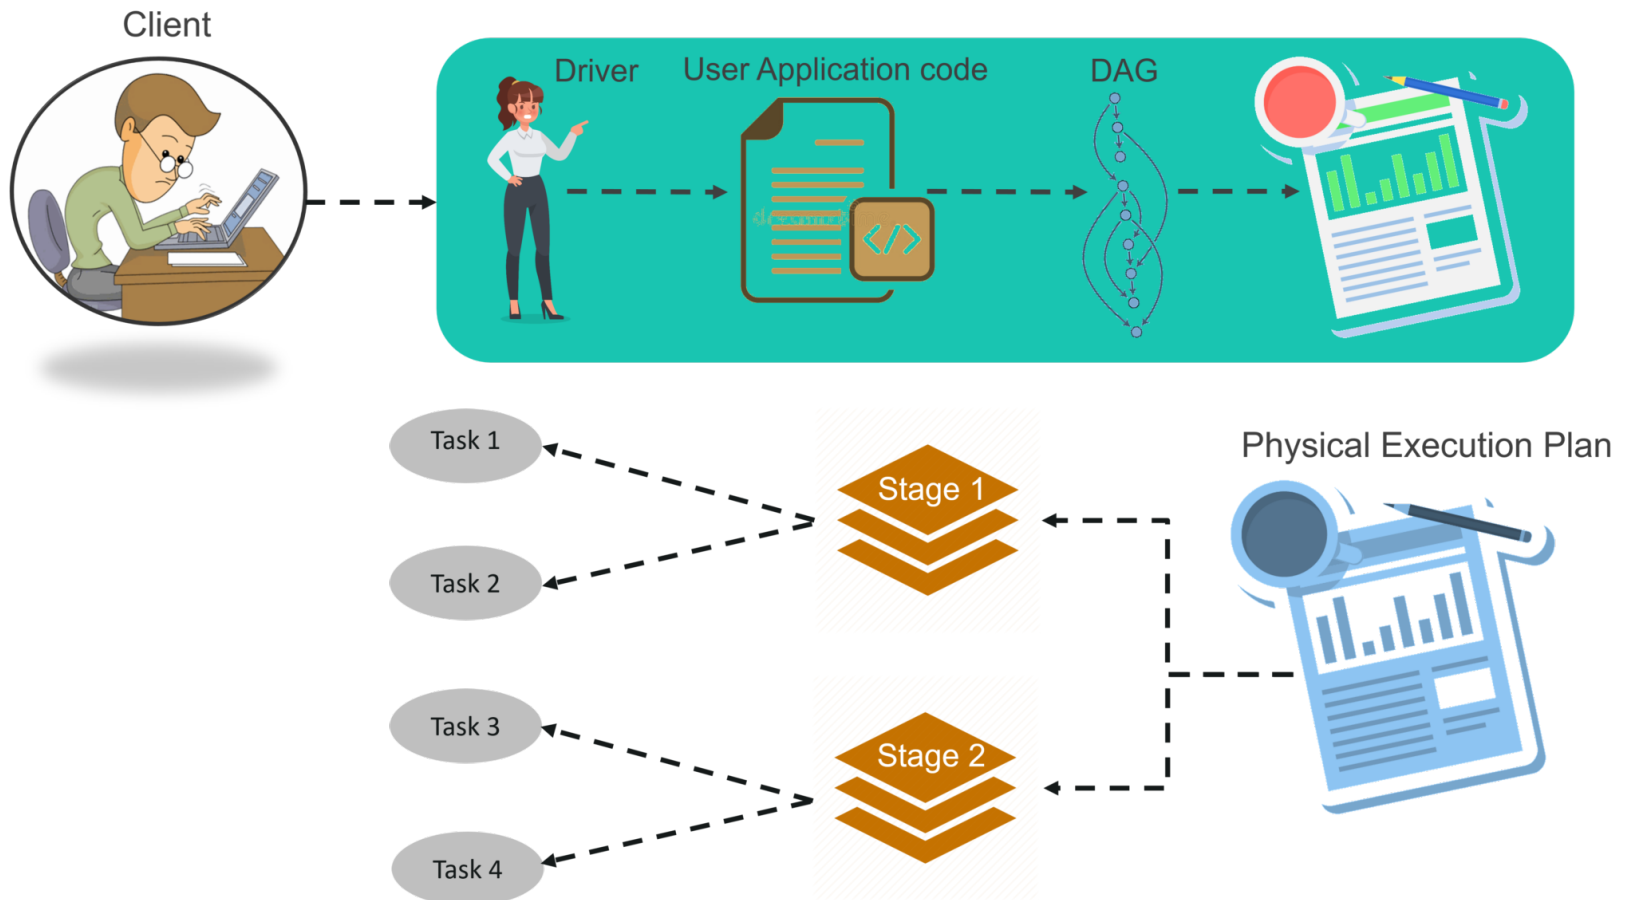
# Spark architecture

- The previous slides presented the key programming abstractions ...
- ... and showed how some operations can influence performance
  - Shuffle operations
- We now present the Spark architecture in more details

# Spark architecture

- When a driver program submits a job to the Spark Context
  1. The Spark context extracts a DAG of operators
  2. The logical DAG is transformed into a physical execution plan
    - Multiple *stages*: each stage contains a sequence of operations with no intermediate data shuffle
  3. Each stage is split into tasks (one for each partition)
  4. Tasks are scheduled on the cluster
    - Where / close to the data they consume
    - Taking into account the dependencies between tasks

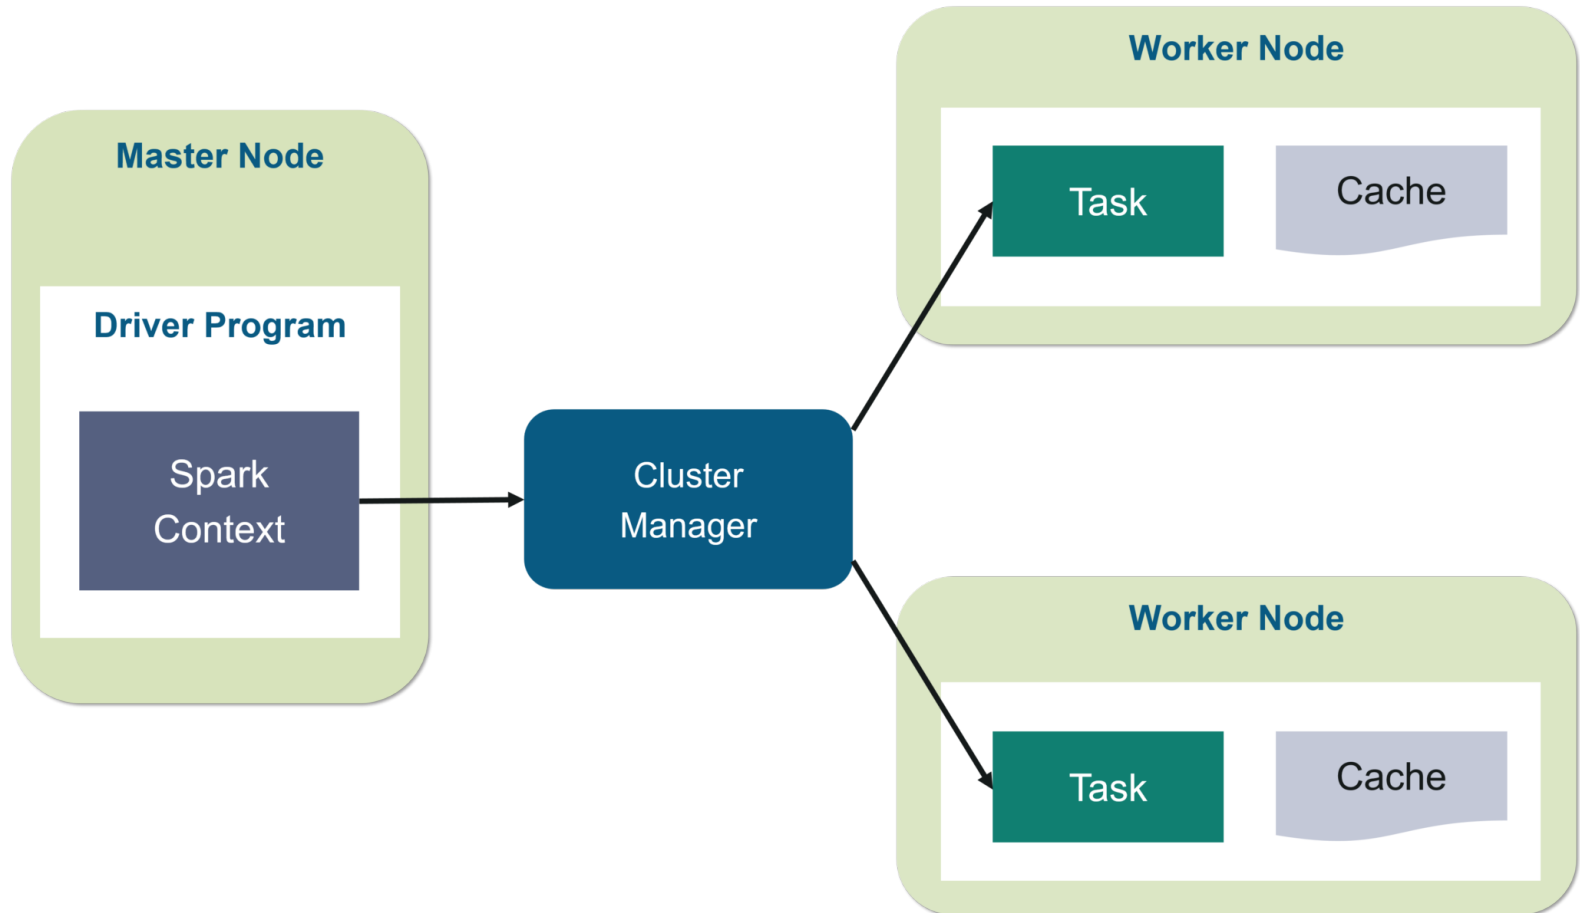
# Spark architecture



# Spark architecture

- When persisted, RDDs are stored in the cache of worker nodes
  - Depending on the specified level, it is cached in memory, on disk, or both
  - Cached in serialized, un-serialized form, or both
- RDDs are partitioned and distributed across workers according to the specified key
- Tasks are scheduled where the data they consume is located

# Spark architecture



# Shared Variables

# Shared variables

- After diving into some architectural details, we can go back to the programming primitives
- As we said, Spark operations consume data and produce new data, ...
- ... they do not operate on shared state
- Spark offers some limited support to shared variables, which can be useful in some scenarios
  - Broadcast variables
  - Accumulators



# Broadcast variables

- If a function passed to Spark accesses a variable ...
- ... it works on a separate copy of that variable in each process (when executed in cluster cluster)
  - Warning: this might not be discovered when testing in local mode!!!
- Broadcast variables enable to keep read-only variables cached in each machine
  - They are used to give to each node a copy of a dataset that everybody needs to read in an efficient manner

# Broadcast variables

- Spark automatically broadcasts the common data within each stage
- Data is cached in serialized form and deserialized before running a task

# Broadcast variables

- Broadcast variables are created from a variable `v` by calling `SparkContext.broadcast(v)`
  - The broadcast variable is a wrapper around `v`
  - The value can be accessed by calling the `value` method

```
val broadcastVar = sc.broadcast(v)  
broadcastVar.value
```

- The original variable should not be modified after it is broadcast ...
- ... to ensure that all the nodes get the same value of the variable when it is delivered

# Broadcast variables

- In some cases, using broadcast variables can be better than accessing an RDD
  - For example, if several operators need to access a static dictionary (that never changes)
- The broadcast variable is stored in the cache of each node in non-serialized form
  - If the dictionary is implemented as a hash map, it is possible to directly retrieve values by key in constant time

# Broadcast variables

- In general, use broadcast variables when you have data that is
  - Not “too large”
  - Shared across multiple operators
    - Not partitioned
  - Read-only

# Accumulators

- Accumulators are variables that can only be modified using associative and commutative operations
  - Example: counters, sums, ...
  - Can be easily supported in parallel
- Accumulators can have an associated name
  - In this case they are displayed in the Web UI
  - They are useful to keep track of tasks

# Accumulators

- Numeric accumulators can be created with
  - `SparkContext.longAccumulator()`
  - `SparkContext.doubleAccumulator()`
- Tasks can add using the `add()` method
- Only the driver can read the accumulator's value using the `value` method

# Accumulators

```
val acc = sc.longAccumulator("Name")
```

```
sc.parallelize(Array(1, 2, 3, 4))  
  .foreach(x => acc.add(x))
```

```
accum.value
```



# Accumulators

- Developers can create custom accumulators by
  - Inheriting from the `AccumulatorV2` class
  - Overriding `reset()` to reset the accumulator to the initial empty / zero value
  - Overriding `add()` to add a new value
  - Overriding `merge()` to merge two values (partial results)
- To register an accumulator `acc` of a custom type `accType`  
`SparkContext.register(acc, "accName")`

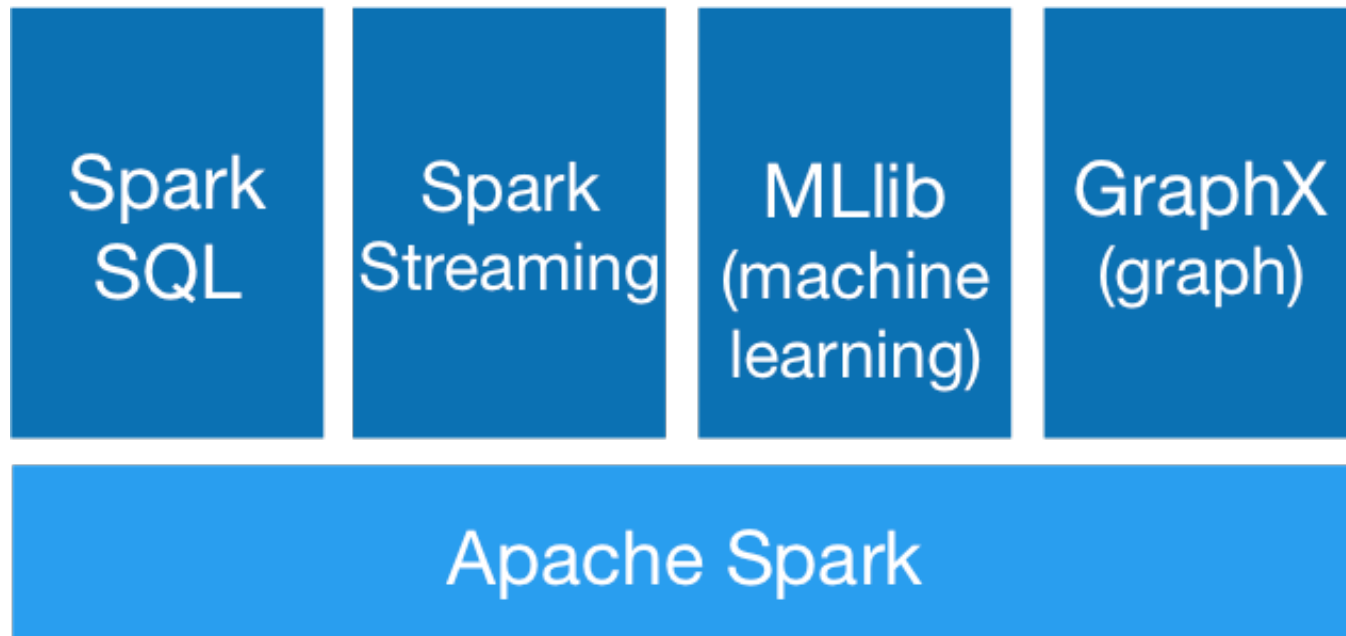
# Accumulators

- Accumulators do not change the lazy evaluation approach of Spark: they are only executed when a transformation is triggered directly or indirectly by an action
- Spark guarantees that accumulators inside actions are updated exactly once
- Warning: Spark does not guarantee that accumulators inside transformations are updated exactly once
  - They can be updated more than once if the task is re-executed

# Spark SQL

SQL, Datasets, and DataFrames

# Spark SQL



# Spark SQL

- Spark SQL is a Spark module for structured data processing
  - Built on top of the core Spark infrastructure
- The interface provides more information about the structure of the data, with multiple advantages
  - Higher-level, declarative API, derived from well known database concepts
  - More opportunities for the engine to optimize the computation

# Spark SQL

- Spark SQL offers different API/languages
  - SQL
  - Dataset API
  - DataFrame
- The same Spark SQL engine is used, independently from the API/language adopted
  - It is possible / easy to switch between different APIs depending on the specific needs of the application

# Spark SQL: SQL API

- The SQL API enables developers to execute standard SQL queries
  - From a programming language
  - From command-line
  - Over JDBC/ODBC
  - ...

# Spark SQL: Dataset and DataFrames

- Dataset is an interface that provides the level of abstraction of RDDs ...
  - Transformations and actions over large collections
- ... with the additional benefits of the Spark SQL optimized execution engine



# Spark SQL: Dataset and DataFrames

- Datasets use a specialized encoder to serialize the objects, processing them, and transmitting over the network
- By knowing the format, Spark can perform many operations without deserializing the object!
  - Filtering, sorting, hashing, ...

# Spark SQL: SparkSession

- The Spark SQL module exposes its functionalities through the SparkSession class
  - Similar to SparkContext for the core Spark API
  - Obtained using a builder

```
val spark = SparkSession
    .builder()
    .appName("App name")
    .master("local")
    .config("option name", "option value")
    .getOrCreate()
```

# Spark SQL: Datasets creation

- Datasets enable the engine to optimize the execution using encoders
  - Differently from a serializer, encoders are code generated dynamically that performs many operations without de-serializing the objects
- In Scala, encoder are available for:
  - Most common types (using implicits)
  - Case classes
- DataFrames can be converted into Datasets by providing a class
  - More on DataFrames, later ...

# Spark SQL: Datasets creation

```
// Dataset from case classes
case class Person(name: String, age: Long)
val ds1 = Seq(Person("Ale", 24)).toDS()

// Dataset from "common" types (integer here)
val ds2 = Seq(1, 2, 3).toDS()

// Dataset from DataFrame
val dataFrame = spark.read.csv("some_path")
val ds3 = dataFrame.as[Person]
```

# Spark SQL: DataFrames

- A DataFrame is a Dataset organized into named columns
- Conceptually equivalent to a relational database
- In Scala, a DataFrame is a Dataset of Row
  - DataFrame is an alias of Dataset[Row]

# Spark SQL: DataFrame creation

- Application can create DataFrames using the methods in `SparkSession.read`
  - Different sources

```
spark.read.json("file")
```

```
spark.read.csv("file")
```

```
spark.read.textFile("file")
```

```
spark.read.jdbc(...)
```

# Spark SQL: DataFrame creation

- DataFrames can also be created from RDDs
  - By inferring the schema using reflection
  - By programmatically defining the schema
- For example, reflection is used with case classes
  - `People` is a case class
  - `peopleRDD` is `RDD[People]`
  - You can obtain a DataFrame using `peopleRDD.toDF()`
  - You can refer to fields (columns) by name, which is inferred by the name of the fields in `People`

# Spark SQL: DataFrame creation

- Instead, it is possible to programmatically define the schema of a DataFrame when it is not known upfront
  - For instance, if it is loaded from a file

```
val mySchema = StructType(Array(  
    StructField("name", StringType, true),  
    StructField("surname", StringType, true),  
    StructField("age", IntegerType, true),  
))
```

```
val df = spark.createDataFrame(myRDD, mySchema)
```



# Spark SQL: DataFrame

- It is possible to print the content of a DataFrame on the standard output

`df.show()`

- Textual representation of the table, ...
- ... very useful for testing and debugging
  - Display the results of transformation on a small dataset

# Spark SQL: DataFrame operations

- DataFrames include additional information on their content (e.g., names of columns)
  - Enables more declarative processing
  - Implicit conversions and simplified column access (using the \$ notation) improve readability

```
import spark.implicits._
```

```
df.filter($"age" > 18)  
  .select($"name", $"salary" + 10)
```

# Spark SQL: DataFrame operations

- SparkSQL also accepts SQL queries as strings

```
employeeDataFrame
```

```
.createOrReplaceTempView("employee")
```


```
val resultDF = spark.sql  
(  
    "    SELECT *  
      FROM employee  
     WHERE salary < 100  ")
```

# Spark SQL: DataFrame operations

- Spark SQL supports common aggregations
  - `count()`, `countDistinct()`, `avg()`, `sum()`, `min()`, `max()`
- It enables developers to write and use their own custom aggregations
  - Extending the `UserDefinedAggregateFunction` class
  - Overriding methods to
    - Initialize the aggregate value
    - Update the aggregate value
    - Merge partial results

# Spark Streaming

# Spark SQL



# Spark Streaming

- Spark Streaming is an extension of the core Spark API to process streaming data
- Other processing engines (e.g., Apache Storm, Apache Flink) adopt a streaming architecture
  - Operators are instantiated and deployed
  - Streams of data flow from operator to operator
  - Pro: very low delay
  - Cons: dynamic adaptation (e.g., dynamic scalability) more difficult
    - Since operators are pre-deployed

# Spark Streaming

- Spark Streaming adopts a different “micro-batch” approach:
  1. It splits the input streams into small batches, ...
  2. ... which are processed by the Spark engine ...
  3. ... to generate the final stream of results in batches



# Spark Streaming



- Pro: dynamic adaptation is easier (scheduling decisions can change over time)
- Cons: higher processing delay

# Spark Streaming API

- Spark Streaming's main abstraction is the *discretized stream (DStream)*
- Internally, a DStream is represented as a sequence of RDDs
- DStreams provide operations to transform the RDDs in the sequence
  - Also provides stateful operations that preserve internal state across invocations

# Spark Streaming Example

- As an example, consider again the word count application
- For the sake of simplicity, we read streaming data from a TCP socket
  - Production environments typically adopt data sources that can be replayed in the case of failures (e.g., Apache Kafka queues)
- When applied in a streaming context, the count is performed separately on each and every RDD

# Spark Streaming Example

```
val conf = new SparkConf()
    .setMaster("local[2]")
    .setAppName("StreamingWordCount")
val ssc = new StreamingContext(conf, Seconds(1))


val counts = ssc.socketTextStream("localhost", 2345)
    .flatMap(_.split(" "))
    .map(word => (word, 1))
    .reduceByKey(_ + _)

counts.print()

ssc.start()
ssc.awaitTermination()
```


# Spark Streaming API

- A DStream is a sequence of RDDs
  - Each RDD contains data from a certain interval
  - More on time, later



# Spark Streaming API

- Any operation applied on a DStream translates to operations on the underlying RDDs



# Receivers

- Input DStreams, received from external sources such as a socket, are associated with a *Receiver* object
  - It received the data from a source and stores it in the memory of Spark
- A receiver occupies one thread
  - When running in local mode, we need to allocate a number of threads  $n$  that is larger than the number of input streams / receivers (“local[ $n$ ]” as the master URL)
  - When running on a cluster, the number of cores allocated to the Spark Streaming applications must be more than the number of receivers

# Receivers: reliability

- Some sources allow the data to be acknowledged
  - E.g., Apache Kafka
- As a consequence, there are two types of receivers
  - Reliable receivers acknowledge the data to the sources once it has been received and stored in Spark with replication
  - Unreliable receivers do not send acknowledgements to a source



# Fault tolerance semantics

- RDDs are stored on durable storage (replicated file system)
- In the case of failure, since transformations are deterministic, the data in transformed RDDs can always be recomputed by re-executing the transformations over the original RDDs
- With Spark Streaming there is an additional problem
  - Data is received from external sources

# Fault tolerance semantics

- In Spark Streaming there are two types of data that the system needs to recover in the case of failure
  1. Data received and replicated: it survives the failure of a node as a copy exists in other nodes
  2. Data received but not yet replicated: the only way to recover this data is to get it again from the source, if possible

# Fault tolerance semantics

- In general, there are three possible types of guarantees when processing streaming data records
  1. At most once: each record is either processed once or not processed at all
  2. At least once: each record is processed one or more times (duplicates are possible in the case of failures)
  3. Exactly once: each record is processed once and only once

# Fault tolerance semantics

- To ensure at least once semantics, sources can send again all the data that was not acknowledged by the receiver
  - This can lead to duplicates in the case the data was received but the acknowledge lost during a failure
- To ensure exactly once semantics, receiving and acknowledging data must be atomic
  - Implemented using some transactional mechanism
  - The data is acknowledged only when it is saved and replicated
  - Data has associated sequential numbers to enable discarding duplicates

# Fault tolerance semantics

- The same holds for receivers, in the case of external processes
- Receivers must also implement some transactional mechanism
  - Ensure that the data is acknowledged only when it is saved on durable state
  - Use sequential numbers to discard duplicates

# Fault tolerance semantics

- In summary, the actual fault tolerance guarantees depend on the specific data sources and sinks
- For instance, exactly once semantics is guaranteed if the input streams and the results are stored on Kafka queues
  - This is a common scenario in modern (micro service) distributed architectures

# DStreams transformations

- DStreams support many of the transformations available on normal RDDs
  - map, flatMap, filter, union, reduce, ...
- As we have seen in the streaming word count example, these transformations are applied separately on each and every RDD in the Dstream

# DStreams transformations

- Another class of interesting transformations are *stateful* operations
  - When processing an element, they preserve some state that can be subsequently accessed while processing further elements
- We will see three examples of stateful operations
  - `updateStateByKey`
  - `mapWithState`
  - `windows`



# DStreams transformations

- `updateStateByKey` creates a *state* DStream
  - This is used to maintain a key-value store
  - The value is updated by applying a given function on the previous state of the key and the new state of the key

# DStreams transformations

// Definition

```
def updateFunction  
(newValues: Seq[Int], oldValue: Option[Int]):  
Option[Int] = {  
    val newValue = ... //  
    Some(newCount)  
}
```

// Application

```
val stateRDD =  
someRDD.updateStateByKey[Int](updateFunction _)
```

# DStreams transformations

- Spark also enables state to be updated and used as part of a transformation
  - Example: `mapWithState`
- We can use this to change the semantics of the streaming word count application
  - The count is preserved and updated across RDDs

# DStreams transformations

```
val initialRDD =  
sc.sparkContext.emptyRDD[Tuple2[String, Int]]
```

```
val stateMapFunction =  
(word: String, count:  
Option[Int], state: State[Int]) => {  
    val sum = count.getOrElse(0) +  
                state.getOption.getOrElse(0)  
    val output = (word, sum)  
    state.update(sum)  
    output  
}
```


# DStreams transformations

```
val counts = textLines
  .map(_._toLowerCase)
  .flatMap(line => line.split(" "))
  .map(word => (word, 1))
  .mapWithState(StateSpec
    .function(stateMapFunction)
    .initialState(initialRDD)
  )
```

# Windows

- Spark Streaming provides windowed computations, to apply transformations over a sliding window of data
- A window is defined in terms of two parameters
  - Window length: the duration of the window
  - Sliding interval: the interval (rate) at which the window operation is performed
- Note: these two parameters must be multiples of the batch interval of the source DStream

# Windows




# Windows

- Spark Streaming offers several operations to define windows and perform computations over windows
  - `countByWindow`
  - `reduceByWindow`
  - `countByKeyAndWindow`
  - `reduceByKeyAndWindow`
  - ...



# Structured Streaming

# Spark SQL



# Structured Streaming

- Alternative API and programming model w.r.t. Spark Streaming
- Build on the Spark SQL engine
- Core ideas
  - Express streaming computations in the same way as batch computations on static data
  - The engine takes care of continuous and incremental execution to update the final results as new data arrives

# Structured Streaming

- Internally, Structured Streaming queries are processed using the Spark micro-batch approach
  - Same latency as Spark Streaming (hundreds of milliseconds)
  - Same fault tolerance semantics (end-to-end exactly once semantics if sources and sinks enable so)
- Spark 2.3 introduced a new Continuous Processing mode
  - Latency in the order of milliseconds


# Programming model

- The core concepts of this programming model are becoming a standard in stream processing
  - SQL / Table API in Flink
  - Kafka Streams / KSQL
  - ...
- Spark Streaming instead, although widely adopted, exploits a programming model that is very engine-specific
  - Micro-batch approach to enable streaming computations on a pure batch-/scheduling-oriented engine

# Programming model


- Key ideas
  - Consider a live stream as a table that is being continuously appended
  - Express streaming computations as standard batch-like queries on static tables
  - Spark runs the computations *incrementally* on the *unbounded* input table
- See the blog post / lecture “Turning the databased inside out with Apache Samza” by M. Kleppmann

# Programming model



Data stream as an unbounded table

# Programming model



Programming Model for Structured Streaming



# Programming model

- A result table / output can be defined in different modes
    - Depending on the need of the sink
1. Complete mode: returns the entire result table
  2. Append mode: returns only the new rows appended to the result table since the last trigger
  3. Update mode: returns only the rows that were updated in the result table since the last trigger
    - If the computation does not contain aggregations, this is equivalent to Append mode

# Programming model: example

- Let us consider again the classic word count example
- To interact with the Spark SQL engine we first need a `SparkSession`

```
val spark = SparkSession
    .builder
    .appName("StructuredStreamingWordCount")
    .getOrCreate()
```

# Programming model: example

```
val lines = spark.readStream
    .format("socket")
    .option("host", "localhost")
    .option("port", 9999)
    .load()
```

```
val words = lines
    .as[String]
    .flatMap(_.split(" "))
```

```
val wordCounts = words
    .groupBy("value")
    .count()
```

- `lines` represents an unbounded table containing streaming text data
  - One “value” column
  - Each line becomes a row in the table
- We convert the `DataFrame` into a `Dataset of String`
- `wordCounts` is again a `DataFrame` containing the count for each word


# Programming model: example

- We can output the results using a query
  - We show the incremental computation model and the results with different modes

```
val query = wordCounts.  
  writeStream  
  .outputMode("complete")  
  .format("console")  
  .start()
```

```
query.awaitTermination()
```

# Programming model: example



Model of the Quick Example

# Incremental execution

- The engine does not materialize the entire table
- Instead, it *incrementally* updates the results upon receiving a new element from the input source
  - It only keeps the minimum intermediate state required to update the result
- This is possible because Structured Streaming implements a set of standard operators with well known semantics

# Event-time and late data

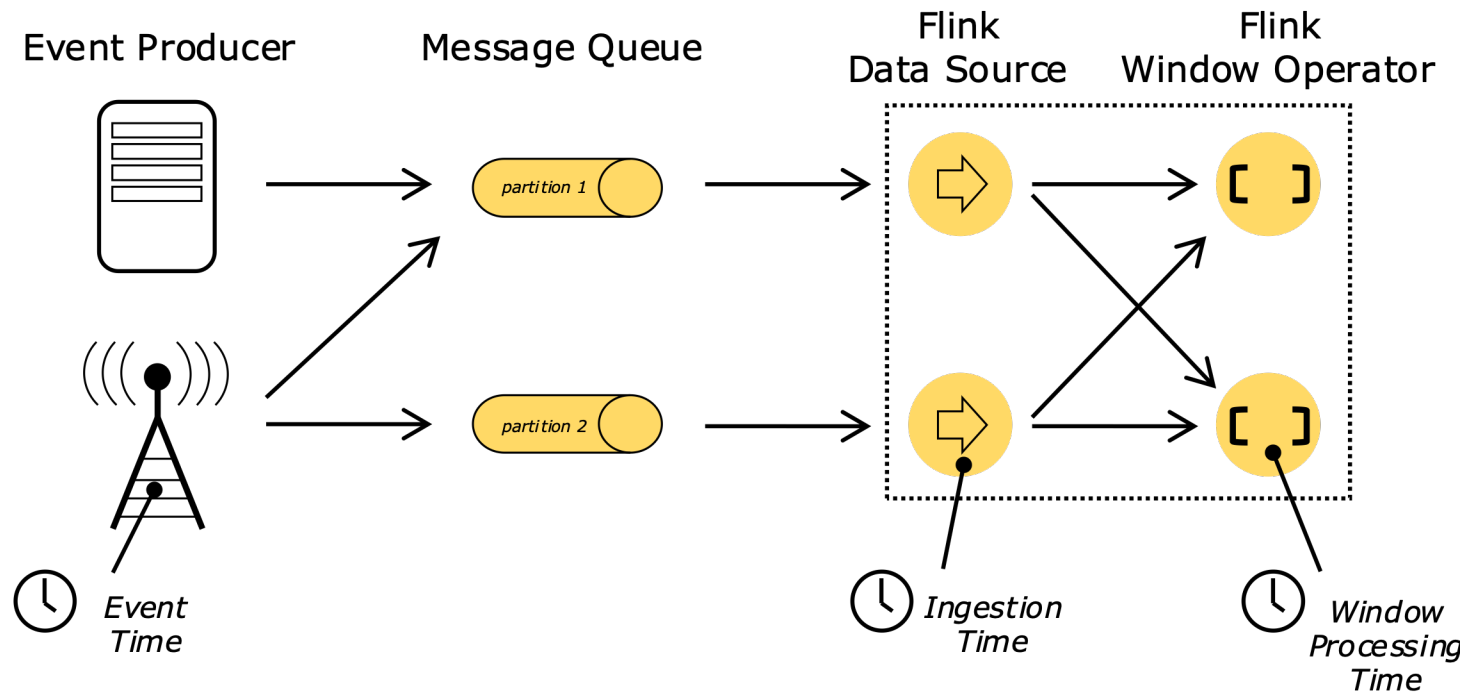
- Some operators rely on time (e.g., windows)
- But what is the meaning of time when running Spark in a distributed environment?
  - Different nodes in the cluster have different clocks
  - Sources and sinks have yet other internal clocks
  - Network communication introduces delays when moving data into/out of the cluster and between the nodes of the cluster

# Event-time and late data

- We can identify three “definitions” of time in stream processing
  1. Event time: is the time attached to a data element by its source
  2. Ingestion time: is the time when a data element first enters the processing cluster
  3. Processing time: is the wall clock time of the processing node



# Event-time and late data



# Event-time and late data

- In most applications, event time is the most significant for the users
  - It is deterministic: in the case of replay, event time does not change and leads to the same results
  - It is set by the application
  - Does not depend on runtime concerns (e.g., load of the processing nodes)

# Event-time and late data

- However, event time is also the most complex to deal with
- In theory, we do not know if old messages are still coming from some source
  - We should wait forever for new messages!
- In practice, we can rely on sources to send information about time
  - *Watermarks*: a watermark with time  $t$  indicates that no further messages older than  $t$  will be received from that source

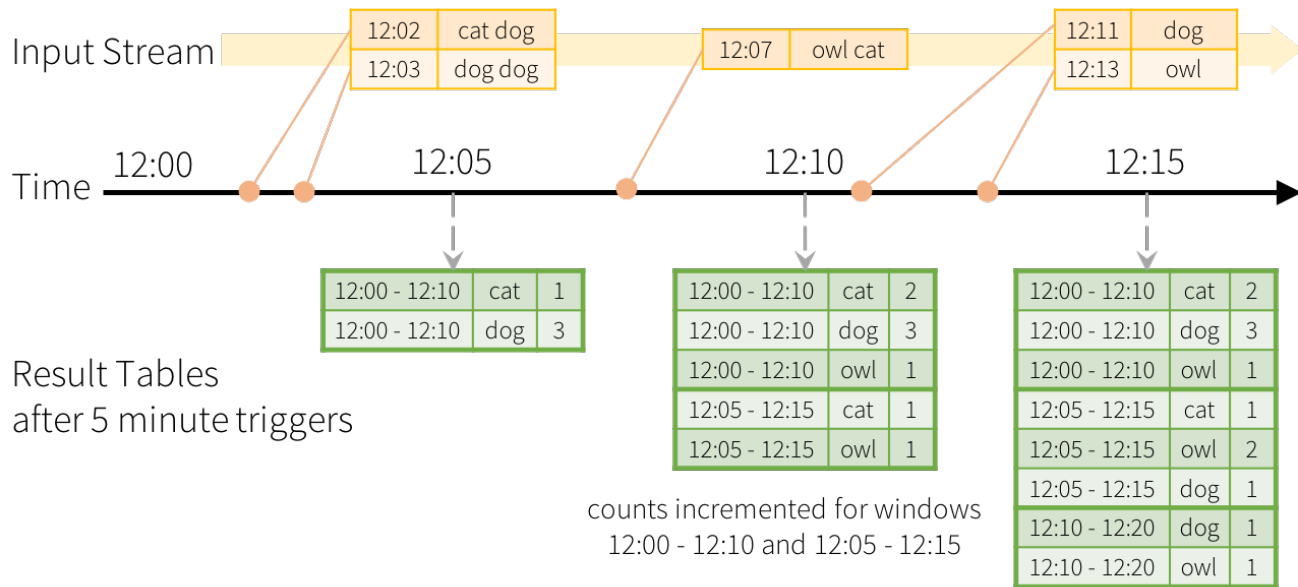
# Event-time and late data

- When receiving a watermark  $t$  from all input sources, we are sure that the results up to time  $t$  are stable
  - We can safely output them
- Structured Streaming takes a different approach
  - Output results are provided immediately
  - They are changed in the case of late arrival
  - To avoid keeping old state forever, watermarks are used to limit the

# Windows

- Aggregations over event-time windows are easy to express
  - Conceptually similar to grouped aggregations
  - Aggregate values are maintained for each window
  - Input data elements can fall into multiple (partially overlapping) time windows
    - They contribute to the value of all the time windows they are part of

# Windows




Windowed Grouped Aggregation  
with 10 min windows, sliding every 5 mins

# Windows and late data

- Window-based grouping is a good example to illustrate how Spark handles late data
- Spark maintain the intermediate state for partial aggregates for a long period of time, such that late data can update aggregates of old windows correctly
- Garbage collection of old intermediate state is handled through watermarking
  - After a threshold, late data is simply discarded

# Windows and late data



Late data handling in  
Windowed Grouped Aggregation



# Joins

- Structured Streaming enables joining streaming datasets with static datasets as well as other streaming datasets
- The result of a streaming join is generated incrementally, similar to the results of streaming aggregations

# Streaming joins: problem

- The problem with joins in streaming data is that a new element in a stream can be potentially joined with any element in another stream
  - We need to store the entire stream, ...
  - ... which grows without bounds over time
- Solutions
  - Specify temporal constraints (time ranges) in the join conditions
    - leftTime BETWEEN rightTime AND rightTime + INTERVAL 30 MINUTES

# Streaming join: solutions

- Define explicit time constraints
  1. Time ranges in the join conditions
    - E.g.  
JOIN ON leftTime  
BETWEEN rightTime AND rightTime + INTERVAL 30 MINUTES
  2. Join on event-time windows
    - E.g.  
JOIN ON leftTimeWindow = rightTimeWindow
- Define watermarks on both input tables
  - The engine knows how late the data is and apply discard policies
  - Similar to streaming aggregations

# Checkpointing

- We already discussed how Spark handles failures
  - End-to-end exactly once semantics
  - By replaying old streaming elements
    - From the receiver
    - From the external sources (e.g., Kafka)
- Problem: in the case of long-running queries, this requires to replay the *entire* stream
  - To restore the set of results
    - E.g., for aggregations over time-based windows
  - To restore the intermediate state for incremental computation

# Checkpointing

- Spark offers the possibility to perform periodic *checkpointing*
  - Store a *snapshot* of the distributed state of the cluster with respect to a query
- Upon failure
  - The state of the cluster is restored from the last valid snapshot
  - The computation restarts from the streaming elements that were not part of the snapshot
    - Replay from receivers or from sources

# Checkpointing

- Developers can specify a checkpoint location for each query
  - The checkpoint location must be in a HDFS-compatible file system

```
val query = streamData
    .writeStream
    .outputMode("complete")
    .option("checkpointLocation", "path/to/HDFS/dir")
    .format("memory")
    .start()
```

# Continuous mode

- Since Spark 2.3, structured streaming also supports a continuous processing mode
  - Similar to other platforms such as Apache Storm and Apache Flink
- Operators are deployed and not scheduled
- They directly exchange messages over the network, ...
- ... rather than accessing the cache of the nodes they work on
  - Memory or disk
- While still experimental, this processing mode enables low delay processing
  - Order of milliseconds

Questions?

